Kevin Stowe, Ph.D.

✉ kevincstowe@gmail.com
📞 +1 517 648 8953
🏠 Chicago, IL, USA
in LinkedIn
🌐 https://kevincstowe.github.io/

About Me

- 🔖 Research scientist and NLP/AI expert.
  - Joint PhD. in Linguistics and Computer Science (2019).
  - 10+ years of experience in natural language processing, deep learning, and linguistics.
  - Projects in natural language generation, LLM exploration, responsible AI, dataset creation and annotation, deepfake detection, crisis informatics, and more.
  - Software development, model training and deployment, evaluation/analysis, academic writing.
  - Focus on responsible AI: detecting and mitigating model bias, model safety, red-teaming, socio-technical solutions.

Skills

Programming	🔖	Python, Java, R, PHP, Javascript, SQL
Toolkits	🔖	HuggingFace, PyTorch/Tensorflow, OpenAI, ScikitLearn, NLTK, Spacy, Gensim
Environments	🔖	Git, AWS (S3, EC2, SageMaker, etc), Linux/Unix/Mac/Windows
Misc.	🔖	Scientific research, writing, and presentation; client communications, team/project leadership, interdisciplinary collaboration
Languages	🔖	Native English, fluent German, beginner Russian

Employment History

2025 – ···· 🔖 **Research Scientist**, Pindrop Security. Atlanta, GA
- Member of Authentication team, working on automating secure communication between customers and call centers via NLP and AI technology.
- Lead research projects in deepfake detection, bias in authentication and security models.
- Implemented improved transformer-based classification of authentication events for customer calls.
- Contributed to Pindrop team on deepfake text detection at NIST 2025 and PAN 2025 competitions.

2024 🔖 **Research Scientist**, Educational Testing Service (ETS), Princeton, NJ

2022 – 2024 🔖 **Associate Research Scientist**
- Automated Content Generation (ACG) team member: built natural language generation systems (LLM and others) for thousands of items over dozens of item types for testing/educational applications.
- Implemented and explored state-of-the-art LLM methodology, with a focus on high-quality evaluation, model fine-tuning, and self-correcting prompting strategies.
- Fairness, bias, and accessibility project leader: explored socio-technical solutions for the responsible generation of educational content.
- Developed and implemented research and infrastructure roadmaps for the generative AI team, incorporating responsible AI, natural language generation, and model evaluation research.
- Member of the Responsible AI working group.

2019 – 2022   **Postdoctoral Researcher**, Ubiquitous Knowledge Processing (UKP), Technical University of Darmstadt, Darmstadt, Germany
● Lead research in metaphor generation, dataset creation for figurative language inferences, and more, yielding top-tier conference publications.
● Developed research in social media during crises, defining foundational research into crowd-sourcing NLP annotation/datasets.
● Instructor for a graduate seminar on Responsible AI focusing on the ethical use of AI, positive societal impacts, and public policy.

2013 – 2019   **Research Assistant**, University of Colorado, Boulder, CO
● NLP expert for the Empowering the Public with Information in Crisis (EPIC) group, focusing on developing positive impacts of NLP for natural and man-made hazards.
● Research in annotation, lexical resources, figurative language, named entity recognition, and computational social science.

Education

2013 – 2019   **Ph.D., University of Colorado**
Joint Degree in Linguistics and Computer Science
Thesis: *Syntactic and semantic improvements to computational metaphor processing*

2009 – 2011   **Master of Arts, Indiana University**
Computational Linguistics Track

2004 – 2009   **Bachelor of Arts, Michigan State University**
Linguistics

Select Publications

1   Stowe, K. *et al.* Identifying Fairness Issues in Automatically Generated Testing Content. *19th Workshop on Innovative Use of NLP for Building Educational Applications.* eprint: 2404.15104 (cs.CL).
🔗 https://arxiv.org/abs/2404.15104 (2024).

2   Pramanick, A., Beck, T., Stowe, K. & Gurevych, I. The challenges of temporal alignment on Twitter during crises. *Findings of Empirical Methods in Natural Language Processing (EMNLP),* 2658–2672.
🔗 https://aclanthology.org/2022.findings-emnlp.195 (Dec. 2022).

3   Stowe, K., Ghosh, D. & Zhao, M. Controlled Language Generation for Language Learning Items. *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing (EMNLP): Industry Track* (eds Li, Y. & Lazaridou, A.) 294–305.
🔗 https://aclanthology.org/2022.emnlp-industry.30 (Dec. 2022).

4   Stowe, K., Utama, P. & Gurevych, I. IMPLI: Investigating NLI Models' Performance on Figurative Language. *Proceedings of the Association for Computational Linguistics (ACL),* 5375–5388.
🔗 https://aclanthology.org/2022.acl-long.369 (May 2022).

5   Stowe, K., Chakrabarty, T., Peng, N., Muresan, S. & Gurevych, I. Metaphor Generation with Conceptual Mappings. *Proceedings of the Association for Computational Linguistics (ACL),* 6724–6736.
🔗 https://aclanthology.org/2021.acl-long.524 (Aug. 2021).